babette: BEAUti 2, BEAST2 and Tracer for R

Richèl J.C. Bilderbeek[1] and Rampal S. Etienne[1]

[1]Groningen Institute for Evolutionary Life Sciences, University of Groningen, Groningen, The Netherlands

May 4, 2018

Summary

1. In the field of phylogenetics, BEAST2 is one of the most widely used software tools. It comes with the graphical user interfaces BEAUti 2, DensiTree and Tracer, to create BEAST2 configuration files and to interpret BEAST2's output files. However, when many different alignments or model setups are required, a workflow of graphical user interfaces is cumbersome.

2. Here, we present a free, libre and open-source package, `babette`: 'BEAUti 2, BEAST2 and Tracer for R', for the R programming language. `babette` creates BEAST2 input files, runs BEAST2 and parses its results, all from an R function call.

3. We describe `babette`'s usage and the novel functionality it provides compared to the original tools and we give some examples.

4. As `babette` is designed to be of high quality and extendable, we conclude by describing the further development of the package.

Keywords: computational biology, evolution, phylogenetics, BEAST2, R

1

1   Introduction

Phylogenies are commonly used to explore evolutionary hypotheses. Not only can phylogenies show us how species (or other evolutionary units) are related to each other, but we also estimate relevant parameters such as extinction and speciation rates. There are many phylogenetics tools available to obtain an estimate of the phylogeny of a given set of species. BEAST2 (Bouckaert *et al.* 2014) is one of the most widely used ones. It uses a Bayesian statistical framework to estimate the joint posterior distribution of estimated phylogenies and model parameters, from one or more DNA, RNA or amino acid alignments (see figure 1 for an overview of the workflow).

It has a graphical and a command-line interface, that both need a configuration file containing alignments and model parameters. BEAST2 is bundled with BEAUti 2 (Drummond *et al.* 2012) ('BEAUti' from now on), a desktop application to create a BEAST2 configuration file. BEAUti has a user-friendly graphical user interface, with helpful default settings. As such, BEAUti is an attractive alternative to manual and error-prone editing of BEAST2 configuration files.

However, BEAUti cannot be called from a command-line script. This implies that when the user wants to explore the consequences of various settings, this must be done manually. This is the managable workflow when using a few alignments and doing a superficial analysis of sensitivity of the reconstructed tree to model settings. For exploring many trees (for instance from simulations) and for more thorough sensitivity analysis, one would like to loop through multiple (simulated) alignments, nucleotide substitution models, clock models and tree priors. One such tool to replace BEAUti is `BEASTmasteR` (Matzke 2015), which focuses on morphological traits and tip-dating, but also supports DNA data. `BEASTmasteR`, however, requires hundreds of lines of R code to setup the

2

BEAST2 model configuration and a Microsoft Excel file to specify alignment files.

BEAST2 is also associated with Tracer (Rambaut & Drummond 2007) and DensiTree (Bouckaert & Heled 2014). Both are desktop applications to analyze the output of BEAST2, each with a user-friendly graphical user interface. Tracer's purpose is to analyze the parameter estimates generated from a BEAST2 run. It shows, among others, the effective sample size (ESS) and time series ('the trace', hence the name) of each variable in the MCMC run. Both ESS and trace are needed to assess the strength of the inference. DensiTree visualizes the phylogenies of a BEAST2 posterior, with many options to improve the simultaneous display of many phylogenies.

However, for exploring the output of many BEAST2 runs, one would like a script to collect all parameters' ESSes, parameter traces and posterior phylogenies. There is no single package that offers a complete solution, but examples of R packages that offer a partial solution are rBEAST (Ratmann 2015) and RBeast (Faria & Suchard 2015). RBeast provides some plotting options and parsing of BEAST2 output files, but the plotting functions are too specific for general use, whilst the functions for parsing call those of `babette`. rBEAST was developed to test a particular biological hypothesis (Ratmann *et al.* 2016), and hence was not designed for general use.

Here, we present `babette`: BEAUti 2, BEAST2 and Tracer for R, which creates BEAST2 (v.2.4.7) configuration files, runs BEAST2, and analyzes its results, all from an R function call. This will save time, tedious mouse clicking and reduces the chances of errors in such repetitive actions. The interface of `babette` mimics the tools it is based on. This familiarity helps both beginner and experienced BEAST2 users to make the step from those tools to `babette`. `babette` enables the creation of a single-script pipeline from sequence alignments

to posterior analysis in R.

2   Description

babette is written in the R programming language (R Core Team 2013) and enables the full BEAST2 workflow from an R function call, in a similar way to what BEAUti, DensiTree and Tracer do. babette's main function is run, which configures BEAST2, runs it and parses its output. run needs at least the name of a FASTA file containing a DNA alignment. The default settings for the other arguments of run are identical to BEAUti's and BEAST2's default settings. Per alignment, a site model, clock model and tree prior can be chosen. Multiple alignments can be used, each with its own (unlinked) site model, clock model and tree prior.

babette currently has 108 exported functions to set up a BEAST2 configuration file. babette can currently handle a majority of BEAUti use cases. Because of BEAUti's high number of plugins, babette uses a software architecture that is designed to be extended. Furthermore, babette has 13 exported functions to run and help run BEAST2. One function is used to run BEAST2, another one installs BEAST2 to a default location. Finally, babette has 21 exported function to parse the BEAST2 output files and analyze the created posterior. babette gives the same ESSes and summary statistics as Tracer. The data is formatted such that it can easily be visualized using ggplot2 (for a trace, similar to Tracer) or phangorn (Schliep 2011) (for the phylogenies in a posterior, similar to DensiTree).

Currently, babette does not replace all functionality in BEAUti, BEAST2 and their many plug-ins. This is reasonable, as these themselves also change in time. babette provides the majority of default tree priors and works only on DNA data. babette supports the most important command-line arguments

4

of BEAST2, provides the core Tracer analysis options, and has the most basic subset of plotting options of DensiTree. Up until now, the `babette` features implemented are those requested by users. Further extension of `babette` will be based on future user requests.

3   Usage

`babette` can be installed easily with the `devtools` package (Wickham & Chang 2016):

```
library ( devtools )
install_github ( " richelbilderbeek / babette " )
```

To start using `babette`, load its functions in the global namespace first:

```
library ( babette )
```

For most convenience, let `babette` install BEAST2 to a default location:

```
install_beast2 ()
```

BEAUti, and likewise `babette`, needs at least a FASTA filename to produce a BEAST2 configuration file. In BEAUti, this is achieved by loading a FASTA file, then saving an output file using a common save file dialog. After this, BEAST2 needs to be applied to the created configuration file. It creates multiple files storing the posterior. These output files must be parsed by either Tracer or DensiTree. In `babette`, all this is achieved by:

```
out <- run ( fasta_filenames = " anthus_aco . fas " )
```

This code will create a (temporary) BEAST2 configuration file, from the FASTA file with name `anthus_aco.fas` (which is supplied with the package, from (Van Els & Norambuena 2018)), using the same default settings as BEAUti, which are, among others, a Jukes-Cantor site model, a strict clock, and a Yule

birth tree prior. `babette` will then execute BEAST2 using that file, and parses the output. The returned data structure, named `out`, is a list of parameter estimates (called `estimates`), posterior phylogenies (called `anthus_aco_trees`, named after the alignment's name) and MCMC operator performance (`operators`). An example of using a different site model, clock model and tree prior is:

```
out <- run(
  fasta_filenames = "anthus_aco.fas",
  site_models = create_hky_site_model(),
  clock_models = create_rln_clock_model(),
  tree_priors = create_bd_tree_prior()
)
```

This code uses an HKY site model, a relaxed log-normal clock model and a birth-death tree prior, each with their default settings in BEAUti. Table 1 shows an overview of all functions to create site models, clock models and tree priors. Note that the arguments' names `site_models`, `clock_models` and `tree_priors` are plural, as each of these can be (a list of) one or more elements. Each of these arguments must have the same number of elements, so that each alignment has its own site model, clock model and tree prior. An example of two alignments, each with its own site model, is:

```
out <- run(
  fasta_filenames = c(
    "anthus_aco.fas",
    "anthus_nd2.fas"
  ),
  site_models = list(
    create_tn93_site_model(),
    create_gtr_site_model()
  )
```

6

```
155  )
```

babette also uses the same default prior distributions as BEAUti for each of the site models, clock models and tree priors. For example, by default, a Yule tree prior assumes that the birth rate follows a uniform distribution, from minus infinity to plus infinity. One may prefer a different ddistribution instead. Here is an example how to specify an exponential distribution for the birth rate in a Yule tree prior in babette:

```
162  out <- run(
163    fasta_filenames = "anthus_aco.fas",
164    tree_priors = create_yule_tree_prior(
165      birth_rate_distr = create_exp_distr()
166    )
167  )
```

In this same example, one may specify the initial shape parameters of the exponential distribution. In BEAST2's implementation, an exponential distribution has one shape parameter: its mean, which can be set to any value with BEAUti. Within babette, to set the mean value of the exponential distribution to a fixed (non-estimated) value, do:

```
173  out <- run(
174    fasta_filenames = "anthus_aco.fas",
175    tree_priors = create_yule_tree_prior(
176      birth_rate_distr = create_exp_distr(
177        mean = create_mean_param(
178          value = 1.0,
179          estimate = FALSE
180        )
181      )
182    )
```

```
183  )
```

NOTE: suggest to remove this paragraph until node dating Our initial motivation to create `babette` was that we wanted to fix the crown age (the time of first divergence) of a phylogeny. BEAUti assumes that a phylogeny has a crown age that needs to be jointly estimated with the phylogeny and other parameters. It does not allow for fixing the crown age. Without `babette`, one needs to manually edit the BEAST2 configuration file, to remove all tree operators that change the height of the root node, which is tedious and prone to errors. Fixing the crown ages is useful when it is unnecessary to accurately model the uncertainty of absolute divergence times. This is how to specify a fixed crown age with `babette`:

```
194  out <- run(
195    fasta_filenames = "anthus_aco.fas",
196    posterior_crown_age = 15
197  )
```

`babette` also supports node dating. Like BEAUti, one can specify Most Recent Common Ancestor ('MRCA') priors. An MRCA prior allows to specify taxa having a common ancestor, including a distribution for the date of that ancestor. An application of such a prior is to specify a (close to) fixed crown age, by including all taxa and assuming a tight normal distribution around the desired crown age. Within `babette`, this is achieved as such:

```
204  posterior <- run(
205    fasta_filenames = "anthus_aco.fas",
206    mrca_priors = create_mrca_prior(
207      taxa_names = get_taxa_names("anthus_aco.fas"),
208      alignment_id = get_alignment_id("anthus_aco.fas"),
209      mrca_distr = create_normal_distr(
```

```
210     mean = create_mean_param(value = 15.0, estimate =
211         FALSE),
212     sigma = create_sigma_param(value = 0.01, estimate =
213         FALSE)
214   )
215  )
216 )
```

217 `babette` allows for the same core functionality as Tracer, in which it works on
218 the values of the parameter estimates sampled in the BEAST2 run. This is
219 called the "trace" (hence the name). The start of the trace is usually discarded,
220 as an MCMC algorithm (such as used by BEAST2) first has to converge to
221 its equilibrium. The start of the trace, called the "burn-in", will be removed,
222 because its parameter estimates are not representative. By default, Tracer dis-
223 cards the first 10% of all the parameter estimates. To remove a 20% burn-in
224 from all parameter estimates in `babette`, the following code can be used:

```
225 traces <- remove_burn_ins(
226   traces = out$estimates,
227   burn_in_fraction = 0.2
228 )
```

229 Tracer shows the ESSes of each posterior's variables. These ESSes are important
230 to determine the strength of the inference. As a rule of thumb, an ESS of 200 is
231 acceptable for any parameter estimate. To calculate the effective sample sizes
232 (of all estimated variables) in `babette`:

```
233 esses <- calc_esses(
234   traces = traces,
235   sample_interval = 1000
236 )
```

Tracer displays multiple summary statistics for each estimated variable: the mean and its standard error, standard deviation, variance, median, mode, geometric mean, 95% highest posterior density interval, auto-correlation time and effective sample size. It displays these statistics per variable. In `babette`, these summary statistics are collected for all estimated parameters at once:

```
sum_stats <- calc_summary_stats(
   traces = traces,
   sample_interval = 1000
)
```

`babette` allows for the same functionality as DensiTree. DensiTree displays the phylogenies in a posterior at the same time scale, drawn one over one another, allowing to see the uncertainty in topology and branch lengths. Within the object `out`, the posterior phylogenies are stored as `anthus_aco_trees`, and can be plotted as such:

```
plot_densitree(phylos = out$anthus_aco_trees)
```

Instead of running the full pipeline, `babette` also allows to only create a BEAST2 configuration file. To create a BEAST2 configuration file, with all settings to default, use:

```
create_beast2_input_file(
   input_filenames = babette::get_babette_path("anthus_aco.
      fas"),
   output_filename = "beast2.xml"
)
```

This file can then be loaded and edited by BEAUti, or run by BEAST2. `babette` also allows to only run BEAST2 from a BEAST2 configuration file:

```
run_beast2(
   input_filename = "beast2.xml",
```

```
264    output_log_filename = "run.log",

265    output_trees_filenames = "posterior.trees",

266    output_state_filename = "final.xml.state"

267 )
```

In this example, we specify the names of the desired BEAST2 output files. These output files can then be inspected with other tools, or to be used to continue a BEAST2 run. Also `run` supports to specify these filenames, and these files are put in a temporary folder by default.

4 babette resources

`babette` is free, libre and open source software available at `http://github.com/richelbilderbeek/babette` and is licensed under the GNU General Public License v3.0. `babette` uses the Travis CI (`https://travis-ci.org`) continuous integration service, which is known to significantly increase the number of bugs exposed (Vasilescu *et al.* 2015) and increases the speed at which new features are added (Vasilescu *et al.* 2015). `babette` has a 100% code coverage, which correlates with code quality (Horgan *et al.* 1994; Del Frate *et al.* 1995). `babette` follows Hadley Wickham's style guide (Wickham 2015), which improves software quality (Fang 2001). `babette` depends on multiple packages, which are `ape` (Paradis *et al.* 2004), `beautier` (Bilderbeek 2018b), `beastier` (Bilderbeek 2018a), `devtools` (Wickham & Chang 2016), `geiger` (Harmon *et al.* 2008), `ggplot2` (Wickham 2009), `knitr` (Xie 2017), `phangorn` (Schliep 2011), `rmarkdown` (Allaire *et al.* 2017), `seqinr` (Charif & Lobry 2007), `stringr` (Wickham 2017), `testit` (Xie 2014) and `tracerer` (Bilderbeek 2018c). We tested `babette` to give a clean error message for incorrect input, by calling `babette` one million times with random or random sensible inputs, using the Peregrine high performance computer cluster. The scripts to do so are supplied with

`babette`.

`babette`'s development takes place on GitHub, `https://github.com/richelbilderbeek/babette`, which accommodates collaboration (Perez-Riverol *et al.* 2016) and improves transparency (Gorgolewski & Poldrack 2016). `babette`'s GitHub facilitates feature requests and has guidelines how to do so.

`babette`'s documentation is extensive. All functions are documented in the package's internal documentation. For quick use, each exported function shows a minimal example. For easy exploration, each exported function's documentation links to related functions. Additionally, `babette` has a vignette that demonstrates extensively how to use it. There is documentation on the GitHub to get started, with a dozen examples of BEAUti screenshots with equivalent `babette` code. Finally, `babette` has tutorial videos that can be downloaded or viewed on YouTube, `https://goo.gl/weKaaU`.

5   Citation of babette

Scientists using `babette` in a published paper can cite this article, and/or cite the `babette` package directly. To obtain this citation from within an R script, use:

```
> citation("babette")
```

6   Acknowledgements

Thanks to Yacine Ben Chehida and Paul van Els for supplying their BEAST2 use cases. Thanks again to Paul van Els for sharing his FASTA files for use by this package. Thanks to Leonel Herrera-Alsina, Raphael Scherrer and Giovanni Laudanno for their comments on this package and article. Thanks to Huw Ogilvie and one anonymous reviewer for reviewing this article. Thanks to

12

rOpenSci, and especially Noam Ross and Guangchuang Yu for reviewing the package's source code. We would like to thank the Center for Information Technology of the University of Groningen for their support and for providing access to the Peregrine high performance computing cluster.

7 Authors' contributions

RJCB and RSE conceived the idea for the package. RJCB created and tested the package, and wrote the first draft of the manuscript. RSE contributed substantially to revisions.

References

Allaire, J., Xie, Y., McPherson, J., Luraschi, J., Ushey, K., Atkins, A., Wickham, H., Cheng, J. & Chang, W. (2017) *rmarkdown: Dynamic Documents for R*. R package version 1.8.

Bilderbeek, R.J. (2018a) beastier: BEAST2 from R. `https://github.com/richelbilderbeek/beastier` [Accessed: 2018-03-16].

Bilderbeek, R.J. (2018b) beautier: BEAUti 2 from R. `https://github.com/richelbilderbeek/beautier` [Accessed: 2018-03-16].

Bilderbeek, R.J. (2018c) tracerer: Tracer from R. `https://github.com/richelbilderbeek/tracerer` [Accessed: 2018-03-16].

Bouckaert, R. & Heled, J. (2014) Densitree 2: Seeing trees through the forest. *bioRxiv*, p. 012401.

Bouckaert, R., Heled, J., Kühnert, D., Vaughan, T., Wu, C.H., Xie, D., Suchard, M.A., Rambaut, A. & Drummond, A.J. (2014) Beast 2: a software platform for bayesian evolutionary analysis. *PLoS Comput Biol*, **10**, e1003537.

337 Charif, D. & Lobry, J. (2007) SeqinR 1.0-2: a contributed package to the R
338 project for statistical computing devoted to biological sequences retrieval and
339 analysis. U. Bastolla, M. Porto, H. Roman & M. Vendruscolo, eds., *Struc-*
340 *tural approaches to sequence evolution: Molecules, networks, populations*, Bi-
341 ological and Medical Physics, Biomedical Engineering, pp. 207–232. Springer
342 Verlag, New York. ISBN : 978-3-540-35305-8.

343 Del Frate, F., Garg, P., Mathur, A.P. & Pasquini, A. (1995) On the correlation
344 between code coverage and software reliability. *Software Reliability Engi-*
345 *neering, 1995. Proceedings., Sixth International Symposium on*, pp. 124–132.
346 IEEE.

347 Drummond, A.J., Suchard, M.A., Xie, D. & Rambaut, A. (2012) Bayesian phy-
348 logenetics with beauti and the beast 1.7. *Molecular biology and evolution*, **29**,
349 1969–1973.

350 Fang, X. (2001) Using a coding standard to improve program quality. *Quality*
351 *Software, 2001. Proceedings. Second Asia-Pacific Conference on*, pp. 73–78.
352 IEEE.

353 Faria, N. & Suchard, M.A. (2015) RBeast. `https://github.com/beast-dev/`
354 `RBeast` [Accessed: 2018-03-02].

355 Gorgolewski, K.J. & Poldrack, R. (2016) A practical guide for improving trans-
356 parency and reproducibility in neuroimaging research. *bioRxiv*, p. 039354.

357 Harmon, L., Weir, J., Brock, C., Glor, R. & Challenger, W. (2008) Geiger:
358 investigating evolutionary radiations. *Bioinformatics*, **24**, 129–131.

359 Horgan, J.R., London, S. & Lyu, M.R. (1994) Achieving software quality with
360 testing coverage measures. *Computer*, **27**, 60–69.

Matzke, N.J. (2015) BEASTmasteR: R tools for automated conversion of NEXUS data to BEAST2 XML format, for fossil tip-dating and other uses. `https://github.com/nmatzke/BEASTmasteR` [Accessed: 2018-02-28].

Paradis, E., Claude, J. & Strimmer, K. (2004) APE: analyses of phylogenetics and evolution in R language. *Bioinformatics*, **20**, 289–290.

Perez-Riverol, Y., Gatto, L., Wang, R., Sachsenberg, T., Uszkoreit, J., Leprevost, F., Fufezan, C., Ternent, T., Eglen, S.J., Katz, D.S. *et al.* (2016) Ten simple rules for taking advantage of git and github. *bioRxiv*, p. 048744.

R Core Team (2013) *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria.

Rambaut, A. & Drummond, A.J. (2007) *Tracer v1.4*. Available from http://beast.bio.ed.ac.uk/Tracer.

Ratmann, O. (2015) rBEAST. `https://github.com/olli0601/rBEAST` [Accessed: 2018-03-02].

Ratmann, O., Van Sighem, A., Bezemer, D., Gavryushkina, A., Jurriaans, S., Wensing, A., De Wolf, F., Reiss, P., Fraser, C. *et al.* (2016) Sources of hiv infection among men having sex with men and implications for prevention. *Science translational medicine*, **8**, 320ra2–320ra2.

Schliep, K. (2011) phangorn: phylogenetic analysis in R. *Bioinformatics*, **27**, 592–593.

Van Els, P. & Norambuena, H.V. (2018) A revision of species limits in neotropical pipits anthus based on multilocus genetic and vocal data. *Ibis*.

Vasilescu, B., Yu, Y., Wang, H., Devanbu, P. & Filkov, V. (2015) Quality and productivity outcomes relating to continuous integration in github. *Proceed-*

385 *ings of the 2015 10th Joint Meeting on Foundations of Software Engineering,*
386 pp. 805–816. ACM.

387 Wickham, H. (2009) *ggplot2: elegant graphics for data analysis.* Springer New
388 York.

389 Wickham, H. (2015) *R packages: organize, test, document, and share your code.*
390 O'Reilly Media, Inc.

391 Wickham, H. (2017) *stringr: Simple, Consistent Wrappers for Common String*
392 *Operations.* R package version 1.2.0.

393 Wickham, H. & Chang, W. (2016) *devtools: Tools to Make Developing R Pack-*
394 *ages Easier.* R package version 1.12.0.9000.

395 Xie, Y. (2014) *testit: A Simple Package for Testing R Packages.* R package
396 version 0.4, `http://CRAN.R-project.org/package=testit`.

397 Xie, Y. (2017) *knitr: A General-Purpose Package for Dynamic Report Genera-*
398 *tion in R.* R package version 1.17.

| Name	Description
run	Run BEAST2
create_gtr_site_model	Create a GTR site model
create_hky_site_model	Create an HKY site model
create_jc69_site_model	Create a Jukes-Cantor site model
create_tn93_site_model	Create a TN93 site model
create_rln_clock_model	Create a relaxed log-normal clock model
create_strict_clock_model	Create a strict clock model
create_bd_tree_prior	Create a birth-death tree prior
create_cbs_tree_prior	Create a coalescent Bayesian skyline tree prior
create_ccp_tree_prior	Create a coalescent constant-population tree prior
create_cep_tree_prior	Create a coalescent exponential-population tree prior
create_yule_tree_prior	Create a Yule tree prior
create_beta_distr	Create a beta distribution
create_exp_distr	Create an exponential distribution
create_gamma_distr	Create a gamma distribution
create_inv_gamma_distr	Create an inverse gamma distribution
create_laplace_distr	Create a Laplace distribution
create_log_normal_distr	Create a log-normal distribution
create_normal_distr	Create a normal distribution
create_one_div_x_distr	Create a 1/X distribution
create_poisson_distr	Create a Poisson distribution
create_uniform_distr	Create a uniform distribution

Table 1: babette's main functions
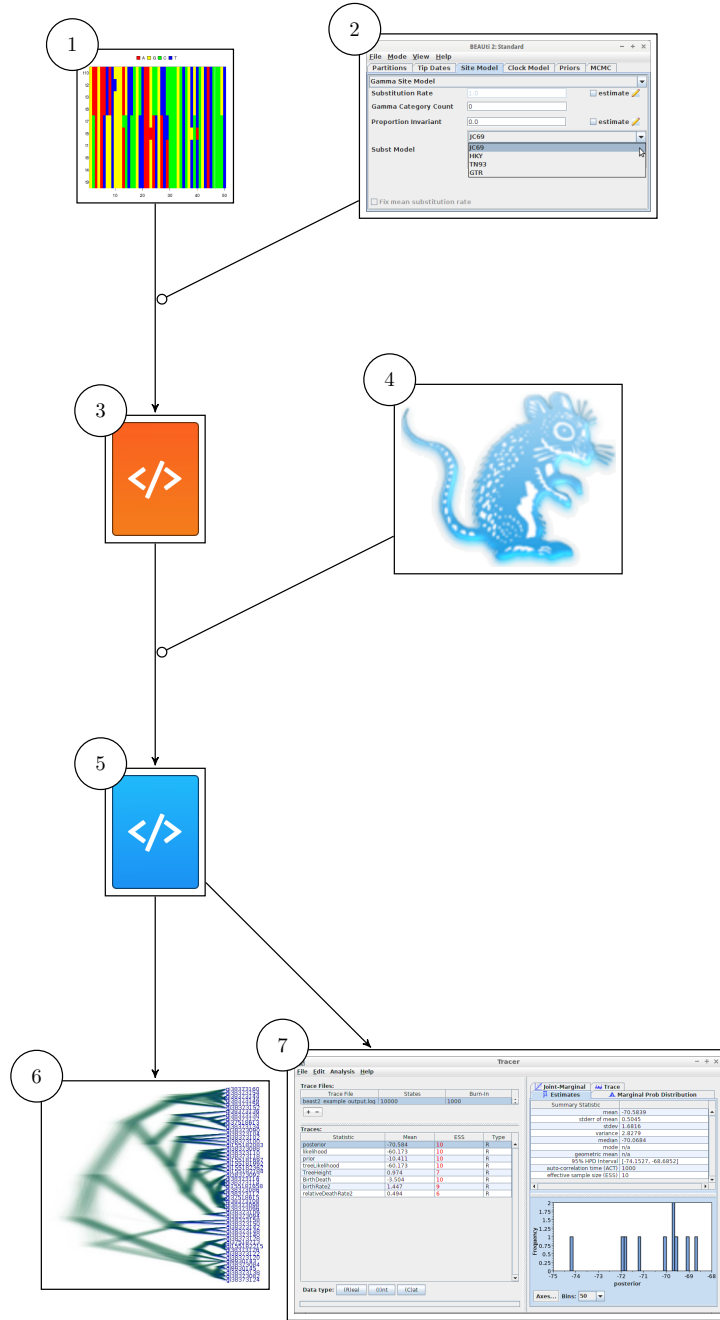
Figure 1: Workflow using GUI tools. From an alignment (1) and BEAUti (2), a BEAST2 configuration file (3) is created. BEAST2 (4) uses that file to infer a posterior, storing it in multiple files (5). These results are visualized using DensiTree (6) and Tracer (7). babette allows for the same workflow, all from an R function call.